

```
In [ ]: cd minbert-final-project
        source setup.sh
        ...
```

```
In [ ]: conda activate nlp_fp
        ...
```

```
In [ ]: def attention(Q, K, V):
        d_k = Q.size(-1)
        scores = torch.matmul(Q, K.transpose(-2, -1)) / math.sqrt(d_k)
        weights = F.softmax(scores, dim=-1)
        output = torch.matmul(weights, V)
        ... return output
        ...
```

```
In [ ]: def multi_head_attention(Q, K, V, num_heads):
        d_model = Q.size(-1)
        head_dim = d_model // num_heads
        Q = Q.view(batch_size, -1, num_heads, head_dim).transpose(1, 2)
        K = K.view(batch_size, -1, num_heads, head_dim).transpose(1, 2)
        V = V.view(batch_size, -1, num_heads, head_dim).transpose(1, 2)

        attention_output = attention(Q, K, V)
        attention_output = attention_output.transpose(1, 2).contiguous().view(batch_size, -1, d_model)
        ... return attention_output
        ...
```

```
In [ ]: def add_norm(x, sublayer):
        ... return LayerNorm(x + sublayer(x))
        ...
```

```
In [ ]: def forward(x):
        x = self.embed(x)
        for layer in self.layers:
            x = layer(x)
        ... return x
        ...
```

```
In [ ]: class BertSentimentClassifier(nn.Module):
        def __init__(self, bert_model):
            super(BertSentimentClassifier, self).__init__()
            self.bert = bert_model
            self.dropout = nn.Dropout(0.1)
            self.classifier = nn.Linear(bert_model.config.hidden_size, num_labels)

        def forward(self, input_ids, attention_mask):
            outputs = self.bert(input_ids, attention_mask=attention_mask)
            pooled_output = outputs[1]
            pooled_output = self.dropout(pooled_output)
            logits = self.classifier(pooled_output)
            ... return logits
            ...
```

```
In [ ]: def train_model(model, train_dataloader, val_dataloader, epochs):
        optimizer = Adam(model.parameters(), lr=2e-5)
        for epoch in range(epochs):
            model.train()
            for batch in train_dataloader:
                optimizer.zero_grad()
                input_ids, attention_mask, labels = batch
                outputs = model(input_ids, attention_mask)
                loss = loss_fn(outputs, labels)
                loss.backward()
                optimizer.step()
            evaluate_model(model, val_dataloader)
        ...
```